

# 7CCSMDLC: Distributed Ledgers & Cryptocurrencies

## *Lecture 3: Mining*



**Peter McBurney**

Professor of Computer Science  
Department of Informatics  
King's College London

Email: [peter.mcburney@kcl.ac.uk](mailto:peter.mcburney@kcl.ac.uk)  
Bush House Central Block North – Office 7.15


2021




# Outline

## Operation of the Bitcoin Blockchain

- Transactions
- Mining and Proof-of-Work
- Reaching Consensus over blocks



# Operation of the Bitcoin Blockchain



# Transactions



# Transaction Outputs

For most transaction, there are two parts:

- An amount of Bitcoin (denominated in satoshis)
- A locking script (an “encumbrance”)
  - The amount is locked unless specific conditions are met

The intended recipient has to provide something to redeem the payment

- Typically they provide their signature (which encodes their private key) and a hash of their public key (their Bitcoin address)
- They may also provide their signature (which encodes their private key) and a hash of a script (a program written in the Bitcoin language Script).
- Some transactions require multiple parties to provide something before the locking script is unlocked.



# UTXO


- **Unspent Transaction Output (UTXO)** is the output of a transaction which may be spent as an input in a subsequent transaction.
- “Sending” a recipient some bitcoin is done by creating some UTXO registered to their address
  - Encumbered to their public key hash or to a script
- All the UTXO of the system is known by every node
  - It is held in a database called the **UTXO set** or **UTXO pool**.
- It is locked to a specific address and may be scattered.
- A wallet will aggregate the UTXO belonging to a single address.



# The 5 Standard Transactions

These are based on what is needed to redeem the payment  
(ie, to satisfy the encumbrance)

- Pay-to-Public-Key-Hash (P2PKH)
  - A hash of a specific public key (a Bitcoin address) is needed to redeem
- Pay-to-Public-Key
  - Mostly used in coinbase transactions
- Multi-sig (multiple-signature)
  - Limited to 15 keys
  - M of N schemes (ie, M signatures of N total signatures are needed, eg 2/3).
- Pay-to-Script-Hash (P2SH)
- Data Output
  - 40 bytes of non-payment data to a Transaction output.



## **Mining & Consensus**





# Four parts of decentralized consensus

Step A Independent verification of each transaction, by every full node

Step B Independent aggregation of those transactions into new blocks by mining nodes, together with demonstrated computation through a Proof-of-Work algorithm

Step C Independent verification of the new blocks by every node and assembly into a chain

Step D Independent selection, by every node, of the chain with the most cumulative computation demonstrated through Proof-of-Work.



# A: Independent verification of transactions

Each node checks against the following list of criteria:

- The transaction's syntax and data structure is correct.
- Neither lists of inputs or outputs are empty.
- The transaction size in bytes is less than `MAX_BLOCK_SIZE`.
- Each output value, as well as the total, is within the allowed range of values
- None of the inputs have `hash=0`, `N=-1` (coinbase transactions should not be relayed)
- `nLocktime` is equal to `INT_MAX`, or `nLocktime` and `nSequence` values are satisfied according to `MedianTimePast`.
- The transaction size in bytes is greater than or equal to 100.
- The number of signature operations (SIGOPS) contained in the transaction is less than the signature operation limit.
- The unlocking script can only push numbers on the stack, and the locking script must match `isStandard` forms.
- A matching transaction in the pool, or in a block in the main branch, must exist.
- For each input, if the referenced output exists in any other transaction in the pool, the transaction is rejected.
- For each input, look in the main branch and the transaction pool to find the referenced output transaction. If the output transaction is missing for any input, this will be an orphan transaction. Add to the orphan transactions pool, if a matching transaction is not already in the pool.
- For each input, if the referenced output transaction is a coinbase output, it must have at least `COINBASE_MATURITY` confirmations.
- For each input, the referenced output must exist and cannot already be spent.
- Using the referenced output transactions to get input values, check that each input value, as well as the sum, are in the allowed range of values (less than 21m coins, more than 0).
- Reject if the sum of input values is less than sum of output values.
- Reject if transaction fee would be too low (`minRelayTxFee`) to get into an empty block.
- The unlocking scripts for each input must validate against the corresponding output locking scripts.



## B: Aggregation TXs into Blocks & Mining



# Mining new bitcoin

- New bitcoin are created during the creation of each block at a fixed and diminishing rate, approx. every 10 minutes.
- Every 210,000 blocks (ca. four years), the currency issuance rate is decreased by 50%
  - 2009-2012: 50 new bitcoin earned per block
  - November 2012: 25 new bitcoin per block
  - July 2016: 12.5 bitcoin per block
  - May 2020: 6.25 bitcoin per block at block 630,000
  - ca. 2137: 1 satoshi per block (block 6,720,000) (99% of all BTC)
  - ca. 2140: After 6.93 million blocks a total of almost 2,099,999,997,690,000 satoshis (almost 21 million bitcoin).
- After that, payment to miners will only be via transaction fees.
- See Block 630,000 here:

<https://blockchair.com/bitcoin/block/630000>

Bitcoin / Block / 629999 — Blockchair

https://blockchair.com/bitcoin/block/629999


Not syncing


**BLOCKCHAIR** Search in 17 blockchains (BTC/ETH/XRP/LTC/BCH/ADA/XLM/BSV/EOS/XMR/XTZ/DASH/ZEC/DOGE/BCHA/XIN/...


Bitcoin Block 629999

Get 100 Spins Buy Crypto Catch your luck Earn Crypto

### General info

  
629998

  
629999

  
630000

Hash: 000000000000000000000000d656be18bb095db1b23bd797266b0ac3ba720b1962b1e

Mined on: 2020-05-11 19:23 (9 months ago) Miner: F2Pool

Coinbase data: [REDACTED] NYTimes 09/Apr/2020 With \$2.3T Injection, Fed's Plan Far Exceeds

Transaction count: 2,481 Fee per kB: 0.00074808 BTC

Witness tx count: 1,528 Fee per kWU: 0.00026740 BTC

Input count: 6,286 Output count: 8,214

Input total: 11,358.76469643 BTC Output total: 11,371.26469643 BTC

Fee total: 1.06876397 BTC Coindays destroyed: 653,689.02

Generation: 12.50000000 BTC Reward: 13.56876397 BTC

### Technical details

Difficulty: 16104807485529

Size: 1,429,136

Weight (weight units): 3,998,681

Stripped size: 856,515

Version: 671080448<sub>10</sub> 27ffe000<sub>16</sub>

Version [bits]: 10011111111111111100000000000002

Median time: 2020-05-11 18:16:54

Merkle root: f4 [REDACTED] 0f

For developers: [API docs](#)

Alternative explorers: [BTC](#) [\[Other\]](#)

[Click to see more](#)

### Transactions included in this block

14

# Reward for mining is new Bitcoin





# The Generation Transaction (Coinbase reward)

- The bitcoin earned by mining are awarded via the first transaction of each new block
  - The Generation (or Coinbase) transaction
- There are no UTXO inputs for these transactions
- Generation transactions do not have an unlocking script (since there is no UTXO). So the field can have arbitrary content:
  - Eg, Satoshi Nakamoto on 03-01-2009 added to the genesis block:

“The Times 03/Jan/2009 Chancellor on brink of second bailout for banks”.

See the Bitcoin Genesis Block (Block #0) here:

<https://blockchair.com/bitcoin/block/0>



17



# Format for the Block Header

Size	Field	Description
4 bytes	Version	Software/protocol version
32 bytes	Previous Block Hash	Reference to previous (parent) block
32 bytes	Merkle Root	Hash of root of merkle tree of the transactions in this block
4 bytes	Timestamp	Creation time of block (seconds from Unix Epoch)
4 bytes	Target	Proof-of-Work algorithm target for this block
4 bytes	Nonce	Counter used for Proof-of-Work algorithm



# Mining problem

- Proof-of-Work is designed to create a hurdle to mining
  - Otherwise, nodes would spin-up multiple sock-puppet nodes to win the reward
  - A form of Sybil attack
- The problems get harder over time
  - To ensure that a new block is created and accepted about every 10 minutes.
- Problem: Find the hash a specified object with a nonce parameter which is less than sum pre-specified total.
  - Problem designed to be hard to do and easy to check.
  - Can only be solved by trial and error.

# Two die example

When throwing two die (dices), how many possible outcomes are there when the total is less than a specified number?

- How many outcomes less than 12 in total
- How many outcomes less than 11 in total
- How many outcomes less than 10 in total
- .....
- How many outcomes less than 3 in total?
- How many outcomes less than 2 in total?
- How many outcomes less than 1 in total?





# Sum of two dice throws

	1	2	3	4	5	6
1	2	3	4	5	6	7
2	3	4	5	6	7	8
3	4	5	6	7	8	9
4	5	6	7	8	9	10
5	6	7	8	9	10	11
6	7	8	9	10	11	12

How many outcomes less than 12 in total: 35 out of 36

How many outcomes less than 11 in total: 33 out of 36

How many outcomes less than 10 in total: 30 out of 36

How many outcomes less than 9 in total: 26 out of 36

.....

How many outcomes less than 3 in total: 1 out of 36




## Example of iterating nonce parameter

I am Satoshi Nakamoto0 => a80a81401765c8eddee25df36728d732...  
I am Satoshi Nakamoto1 => f7bc9a6304a4647bb41241a677b5345f...  
I am Satoshi Nakamoto2 => ea758a8134b115298a1583ffb80ae629...  
I am Satoshi Nakamoto3 => bfa9779618ff072c903d773de30c99bd...  
I am Satoshi Nakamoto4 => bce8564de9a83c18c31944a66bde992f...  
I am Satoshi Nakamoto5 => eb362c3cf3479be0a97a20163589038e...  
I am Satoshi Nakamoto6 => 4a2fd48e3be420d0d28e202360cfbaba...  
I am Satoshi Nakamoto7 => 790b5a1349a5f2b909bf74d0d166b17a...  
I am Satoshi Nakamoto8 => 702c45e5b15aa54b625d68dd947f1597...  
I am Satoshi Nakamoto9 => 7007cf7dd40f5e933cd89fff5b791ff0...  
I am Satoshi Nakamoto10 => c2f38c81992f4614206a21537bd634a...  
I am Satoshi Nakamoto11 => 7045da6ed8a914690f087690e1e8d66...  
I am Satoshi Nakamoto12 => 60f01db30c1a0d4cbce2b4b22e88b9b...  
I am Satoshi Nakamoto13 => 0ebc56d59a34f5082aaef3d66b37a66...  
I am Satoshi Nakamoto14 => 27ead1ca85da66981fd9da01a8c6816...  
I am Satoshi Nakamoto15 => 394809fb809c5f83ce97ab554a2812c...

# Intending miners

- When a new block arrives, miners tackle the next PoW problem
- Meanwhile, they assemble transactions that are not in a block into a candidate block
  - Prioritized by age (how many blocks since the UTXO was recorded) &
  - Size of transaction
- High priority:
  - 1 Bitcoin, aged 1 day
- As new blocks added, unused TXs increase in age
- When miner is restarted, its TX pool is wiped.





## Four parts of decentralized consensus: C & D

Step C: Independent verification of the new blocks by every node and assembly into a chain

Step D: Independent selection, by every node, of the chain with the most cumulative computation demonstrated through Proof-of-Work.

- We can reference blocks by their height (currently about 669,000), or the hash of their header.
  - Block height may not be unique (if there is a fork).
- Block hash is not stored within the block
  - It is calculated by each node as the block is received.






# Validating a new block


Criteria for validation include:

- The block data structure is syntactically valid (ie, format is correct)
- The block header hash is less than the target (enforces Proof-of-Work)
- The block timestamp is less than two hours in the future (allowing for time errors)
- The block size is within acceptable limits
- The first transaction (and only the first) is a coinbase transaction
- All transactions within the block are valid using the transaction checklist for Independent Verification of Transactions.


Blockchain assumes a peer-to-peer (P2P) network  
No one is in control.



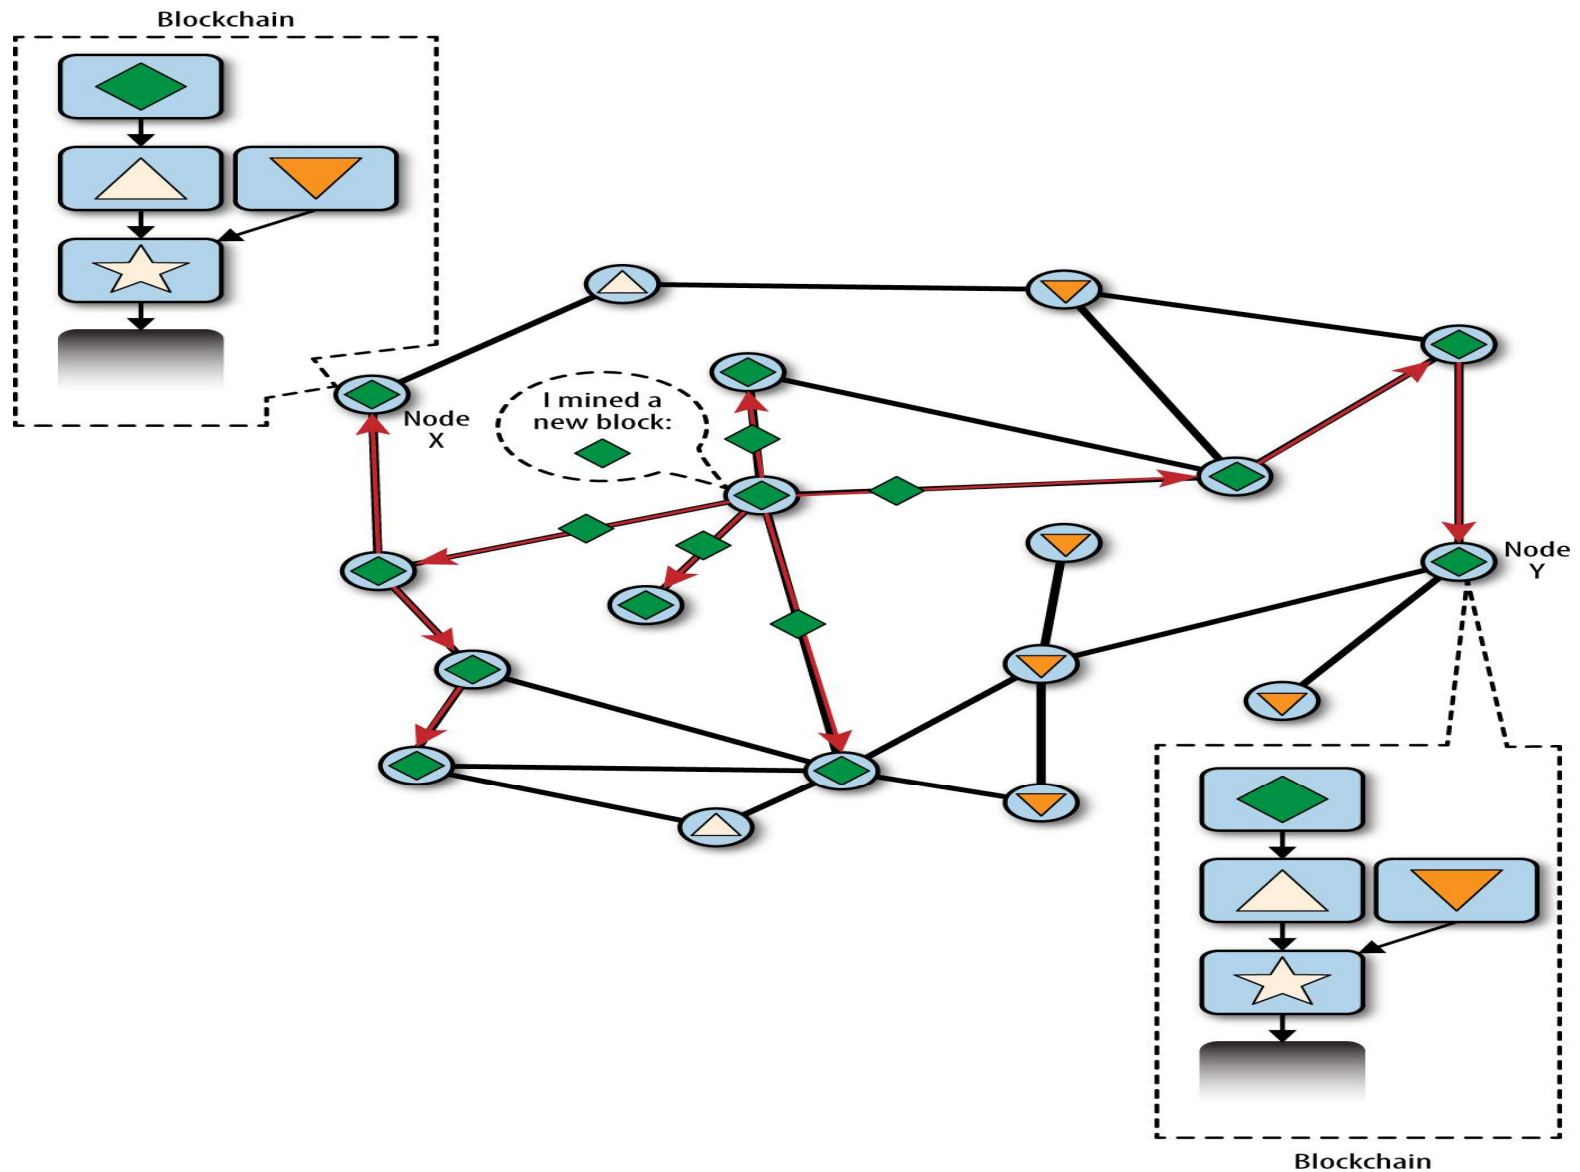
# Nodes mine blocks and propagate them locally



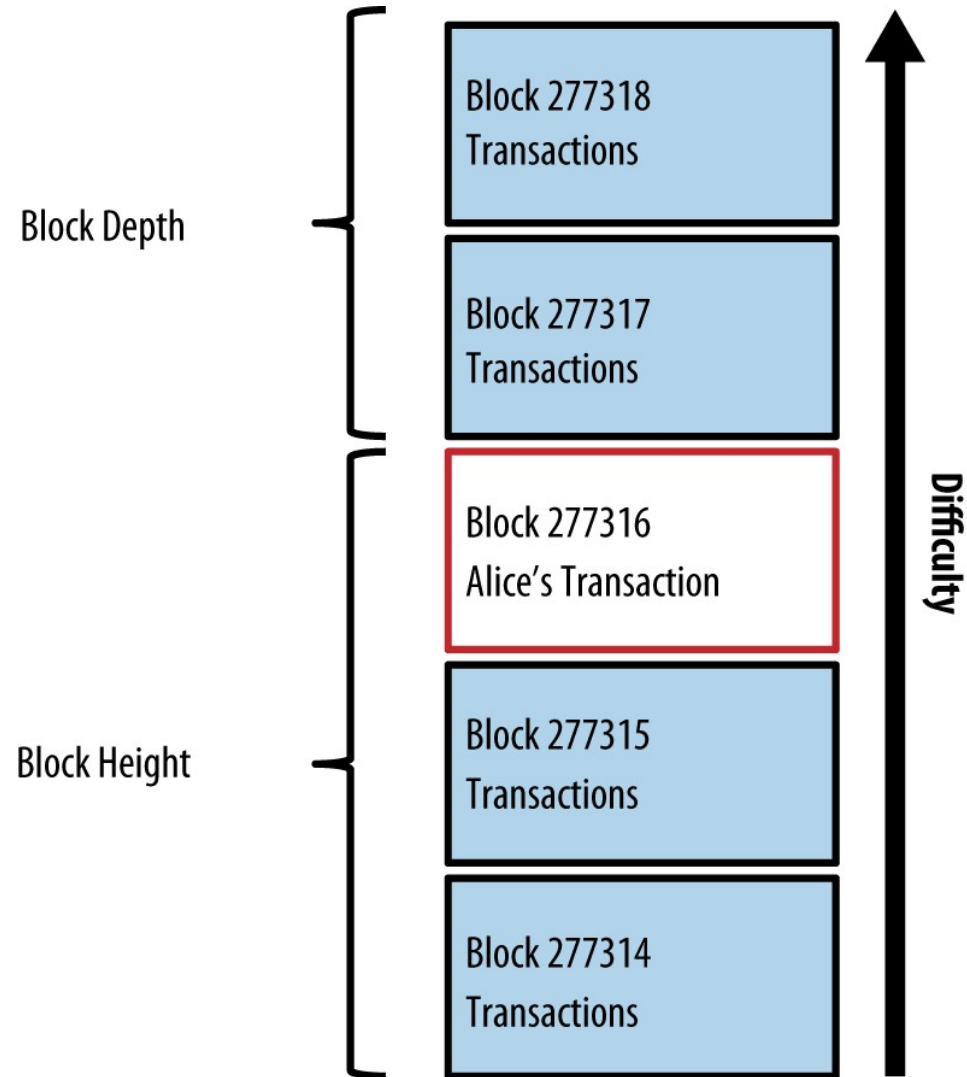
# Competing new blocks from different miners



Which chain is “longer” (contains more work)?



Block height currently is about 669,000



<https://blockchain.info/q/getblockcount>

31



Bitcoin / Block / 668042 — Block

New tab

←

→

↻

🏠

🔒

https://blockchair.com/bitcoin/block/668042

🌟

🔖

Not syncing

⋮

Input total

2,169.27789065 BTC

Output total

2,175.52789065 BTC

Fee total

0.21520995 BTC

Coindays destroyed

5,798.78

Generation

6.25000000 BTC

Reward

6.46520995 BTC

Alternative explorers

BTC

🔍

[Click to see more ↓](#)

Transactions included in this block

Block #	Hash	Inputs #	Outputs #	Coindays destroyed	Output (BTC)	Output (USD)	Transaction fee (BTC)	Transaction fee (USD)	Fee/kB (BTC)	Size (kB)
668042	4f7b	1	3	0.00	6.46520995	194,454.00	0.00000000	0.00	0.00000000	0.247
668042	907a	9	12	1,629.36	1,407.29027588	42,327,100.00	0.00363000	109.18	0.00117704	3.084
668042	f91b	1	4	0.07	4.77249091	143,542.00	0.00054000	16.24	0.00188153	0.287
668042	85a7	2	1	0.00	0.00469994	141.36	0.00043482	13.08	0.00127513	0.341
668042	d44e	1	2	0.00	0.08270061	2,487.39	0.00040320	12.13	0.00161928	0.249
668042	f557	1	2	0.01	4.40783554	132,574.00	0.00039840	11.98	0.00161296	0.247
668042	f95f	1	2	0.03	16.34465500	491,598.00	0.00039840	11.98	0.00161296	0.247
668042	72ea	1	2	0.33	3.75875904	113,052.00	0.00039840	11.98	0.00160645	0.248
668042	77a5	4	1	0.00	0.03894625	1,171.39	0.00084847	25.52	0.00116229	0.730
668042	498e	3	1	0.07	0.03890000	1,170.00	0.00100799	30.32	0.00207833	0.485

[Explore this list \(2,112 more rows\) →](#)

🏠

Type here to search

🔍

📁

📧

📄

📅

📊

🔥

↑

🔌

📶

ENG

13:03

28/01/2021

🔔

32







# How do nodes decide between competing blocks?

- Nodes keep three collections of blocks
  - Those on the main blockchain
  - Those that form branches off the main blockchain
  - Orphan blocks – those without a parent block
- The main chain is the chain with the most cumulative difficulty associated with it
  - Usually the chain with the most blocks
  - If two chains are equal length, then the main chain is the one with most PoW
- Forks usually resolved within 1 block
- 10 minutes for each block time is a compromise between
  - Fast confirmation times & the probability of a fork.

# Eventually consensus is achieved



# Thank you!



[peter.mcburney@kcl.ac.uk](mailto:peter.mcburney@kcl.ac.uk)